Himanshu Yadav

radhekrishna8267@gmail.com linkedin.com/in/himanshuy26/ github.com/Himanshu-2678 Medium

Summary

AI-focused engineer with hands-on experience building and deploying end-to-end ML and NLP systems. Strong in Python, machine learning, deep learning, and transformer-based models, with practical experience integrating LLMs into web apps using Flask. Interested in building scalable, real-world AI systems that actually work in production.

Skills

Programming Languages: Python, SQL
AI/ML: Machine Learning, Deep Learning, NLP, Transformers, Scikit-Learn, TensorFlow, Keras
Libraries: NumPy, Pandas, Matplotlib, Seaborn
Tools & Frameworks: Git, GitHub, Flask, Streamlit, Google Colab
Other Skills: Backend Development (Flask), Model Deployment, Problem Solving

Projects

MeetSync – AI Meeting Minutes Generator | *Live Link* | *Tech: Python, Flask, NLP*

- Built an AI-powered system that converts meeting audio into structured minutes, including discussion topics, decisions, and action items, reducing manual note-taking effort by 60–70%.
- Implemented an end-to-end pipeline combining speech-to-text transcription using the **Deepgram API** with LLM-based reasoning via Google Gemini.
- Developed a lightweight **Flask**-based web application allowing users to upload audio files and download generated meeting summaries.
- Designed session-based handling of summaries to ensure secure, isolated processing for each user.
- Deployed the application on **Render** while optimizing memory usage to operate on a free-tier cloud platform.

Retention Risk Scoring Model | *Project Link* | *Tech: Python, TensorFlow, Keras, Machine Learning*

- Built a binary classification model to predict bank customer churn using demographic features.
- Preprocessed a dataset of **10K+ customer records** using encoding and feature scaling techniques.
- Trained an Artificial Neural Network (ANN) using TensorFlow/Keras for churn risk prediction.
- Achieved **86% test accuracy** and deployed the inference pipeline through an interactive Streamlit application.

Energy Consumption Forecasting | *Project Link* | *Tech: Python, Machine Learning, XGBoost*

- Built a machine learning model to predict the hourly energy consumption for the PJM Interconnection.
- Cleaned and preprocessed data of **167K records** with 13 features from multiple regions.
- Engineered new features such as time of day, hour, day, month, holiday status, etc.
- Applied machine learning models including XGBoost, Decision Trees, and Random Forest.
- Achieved **R^2 score of 0.92** on the test data after hyperparameter tuning to optimize model performance.

Achievements & Certifications

- **AI Azure Virtual Internship: Microsoft & Eduskills with AICTE:** Built a Machine Learning model on a dataset - Algerian Forest Fire. Used Linear Regression to predict fire zones and achieved 98% of accuracy.
- **Hackaccino 3.0 National Hackathon:** Built a Women Safety Platform to predict risk or threat in a given area based on previous crimes on women in that region. Achieved 13/194 position.
- **Data Science and Machine Learning Bootcamp - Udemy**

Education

B.Tech in Computer Science & Engineering **Expected: 2027**
Galgotias University, Gautam Buddha Nagar