Himanshu Yadav

radhekrishna8267@gmail.com | linkedin.com/in/himanshuy26 | github.com/Himanshu-2678 | Blogs: Medium

Summary

Applied Machine Learning Engineer with hands-on experience building and deploying end-to-end ML and NLP systems using Python. Proficient in deep learning, transformer-based models, and integrating large language models into web applications. Interested in building scalable, real-world AI systems that perform reliably in production.

Skills

Programming Languages: Python, SQL
AI/ML: Machine Learning, Deep Learning, Feature Engineering, Model Evaluation, Natural Language Processing, Transformers, Scikit-learn, TensorFlow
Libraries: NumPy, Pandas, Matplotlib, Seaborn
Tools & Frameworks: Git, GitHub, Flask, Streamlit, Google Colab
Other Skills: Backend Development (Flask), Model Deployment, Problem Solving

Projects

MeetSync – AI Meeting Minutes Generator | *Live Demo* | *Tech: Python, Flask, NLP*

- Built an AI-powered system that converts meeting audio into structured minutes, including discussion topics, decisions, and action items, reducing manual note-taking effort by 60–70%.
- Implemented an end-to-end pipeline combining speech-to-text transcription using the Deepgram API with LLM-based reasoning via Google Gemini.
- Developed a lightweight **Flask**-based web application allowing users to upload audio files and download generated meeting summaries.
- Designed session-based handling of summaries to ensure secure and isolated processing for each user.
- Deployed the application on **Render** while optimizing memory usage to operate on a free-tier cloud platform.

Retention Risk Scoring Model | *Project Link* | *Tech: Python, TensorFlow, Keras, Machine Learning*

- Built a binary classification model to predict bank customer churn using demographic features.
- Preprocessed a dataset of more than **10,000 customer records** using encoding and feature scaling techniques.
- Trained an Artificial Neural Network (ANN) using TensorFlow and Keras for churn risk prediction.
- Achieved **86% test accuracy** and deployed the inference pipeline via an interactive Streamlit application.

Energy Consumption Forecasting | *Project Link* | *Tech: Python, Machine Learning, XGBoost*

- Built a machine learning model to predict hourly energy consumption for the PJM Interconnection.
- Cleaned and preprocessed a dataset of **167,000 records** with 13 features from multiple regions.
- Engineered new features such as time of day, hour, day, month, and holiday status.
- Applied machine learning models including XGBoost, Decision Trees, and Random Forest.
- Achieved an **R^2 score of 0.92** on test data after hyperparameter tuning to optimize model performance.

Achievements & Certifications

- **AI Azure Virtual Internship (Microsoft & Eduskills, AICTE):** Built a machine learning model on the *Algerian Forest Fire* dataset using Linear Regression to predict fire zones, achieving **98% accuracy**.
- **Hackaccino 3.0 National Hackathon:** Built a women's safety platform to predict risk levels in a given area based on historical crime data; secured **17th place out of 194 teams**.
- **Data Science and Machine Learning Bootcamp – Udemy**

Education

B.Tech in Computer Science and Engineering **Expected 2027**
Galgotias University, Gautam Buddha Nagar